


Telco Analytics AI System



Over 70% of enterprise data
goes unused for analytics



Project Overview & Business Value

Objective: Enable telecom stakeholders to query business data using natural language.

Solution: A full-stack AI system that translates natural language questions into SQL, executes them on Azure SQL, and visualizes the result.

Business Value:

- Empowers non-technical users to query data easily
- Speeds up telecom business analysis and decision-making
- Bridges the gap between business questions and data insights using LLMs

Tech Stack:

- **SQL generation:** defog/sqlcoder-7b (GGUF quantized version)
- **Backend:** FastAPI + Streamlit
- **Charting:** QuickChart API
- **DB:** Azure SQL (data uploaded via Azure Data Studio)

System Architecture



Project Lifecycle

User Input: Natural Language Query

SQL Generation: Prompt passed to SQLCoder (LLM)

Execution: Query executed on Azure SQL

Post-Processing: DataFrame returned

Charting: Visualization generated

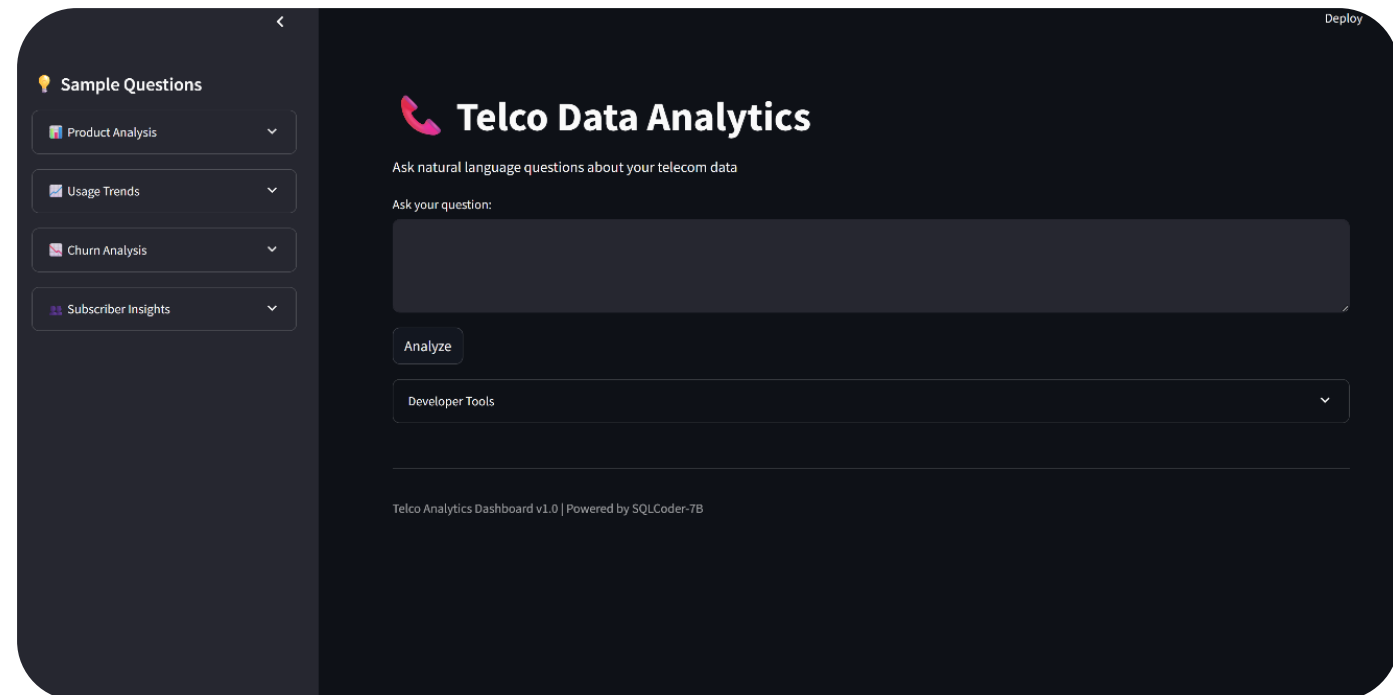
Monitoring: Query performance tracked

Model Used: [defog/sqlcoder-7b](https://huggingface.co/defog/sqlcoder-7b)

Variant: Quantized sqlcoder-7b.Q4_K_M.gguf via llama-cpp for local inference

Input: Natural language telecom-related query

Output: Validated T-SQL queries compatible with Azure SQL



Model Enhancements and Validation Summary

SQL Processing Workflow

1. Input Sanitization

- Remove markdown code blocks (sql)
- Trim whitespace
- Validate query length (> 20 characters)

2. Syntax Correction

- Fix common SQL syntax issues:
 - Malformed or missing aliases
 - Unmatched brackets
 - Stray characters
 - Missing JOIN conditions
 - Inconsistent spacing

Domain-Specific Mapping Rules

Category	Validation	Example
Basic Structure	Starts with SELECT or WITH	SELECT or WITH
Table Reference	[dbo].[table] format	ON clause
JOIN Conditions	Using T-SQL, not PostgreSQL	Blocked join commands
Security	Blocked SQL	

Prompt Engineering Enhancements

- Inject schema context dynamically
- + 22 hand-crafted T-SQL rules
- 5 few-shot examples
- Classify question type to guide generation

Post-Processing Steps

- Convert PostgreSQL outputs to T-SQL

Demo Scenarios



"Show monthly data usage trends for the past year"



"What are the top churn reasons last year?"



"Top 5 most expensive products"

Edge Cases & Error Handling

Case	Behavior
Invalid SQL	Triggers "Execution Error" return
Chart fails	Returns valid SQL/data but reports "Visualization Error"
No results	Returns: No results found
Model fails	Returns: Generation Error

Metrics & Monitoring

Metric	Calculation	Storage
Success Rate	$(\text{Successful Executions} / \text{Total Queries}) \times 100$	session_metrics
Avg Generation Time	$\Sigma(\text{Generation Times}) / \text{Count}$	response_times
Error Frequency	Error Type Counts	common_errors

Conclusion and Future Enhancements

Conclusion

Built a full working MVP of an LLM-powered telecom data analytics platform.
Enables non-technical users to access business insights.

Future Enhancements

- 1) Currently using quantized GGUF due to limited GPU (4GB). Will move to original defog/sqlcoder-7b requiring 6GB+ VRAM.
- 2) Fine-tuned model on telecom schema.
- 3) To save dashboards and query history create user authentication
- 4) Support JOIN-heavy questions